

Scikit-build

**A Build System Generator for
CPython C/C++/Fortran/Cython Extensions**

Speaker: Jean-Christophe Fillion-Robin
Scipy 2018

Who am I ?

- Principal Engineer @ [Kitware](#) in our North Carolina office, lead developer of [3D Slicer](#)
- Maintainer of [scikit-build](#), [cmake](#) and [ninja](#) python packages
- Maintainer of [python-cmake-buildsystem](#)
- Maintainer of [dockcross](#)

Goal of this talk

- Identify requirements for building CPython Binary Extensions
- Outline limitations of current approach
- Introduce our solution

Outline

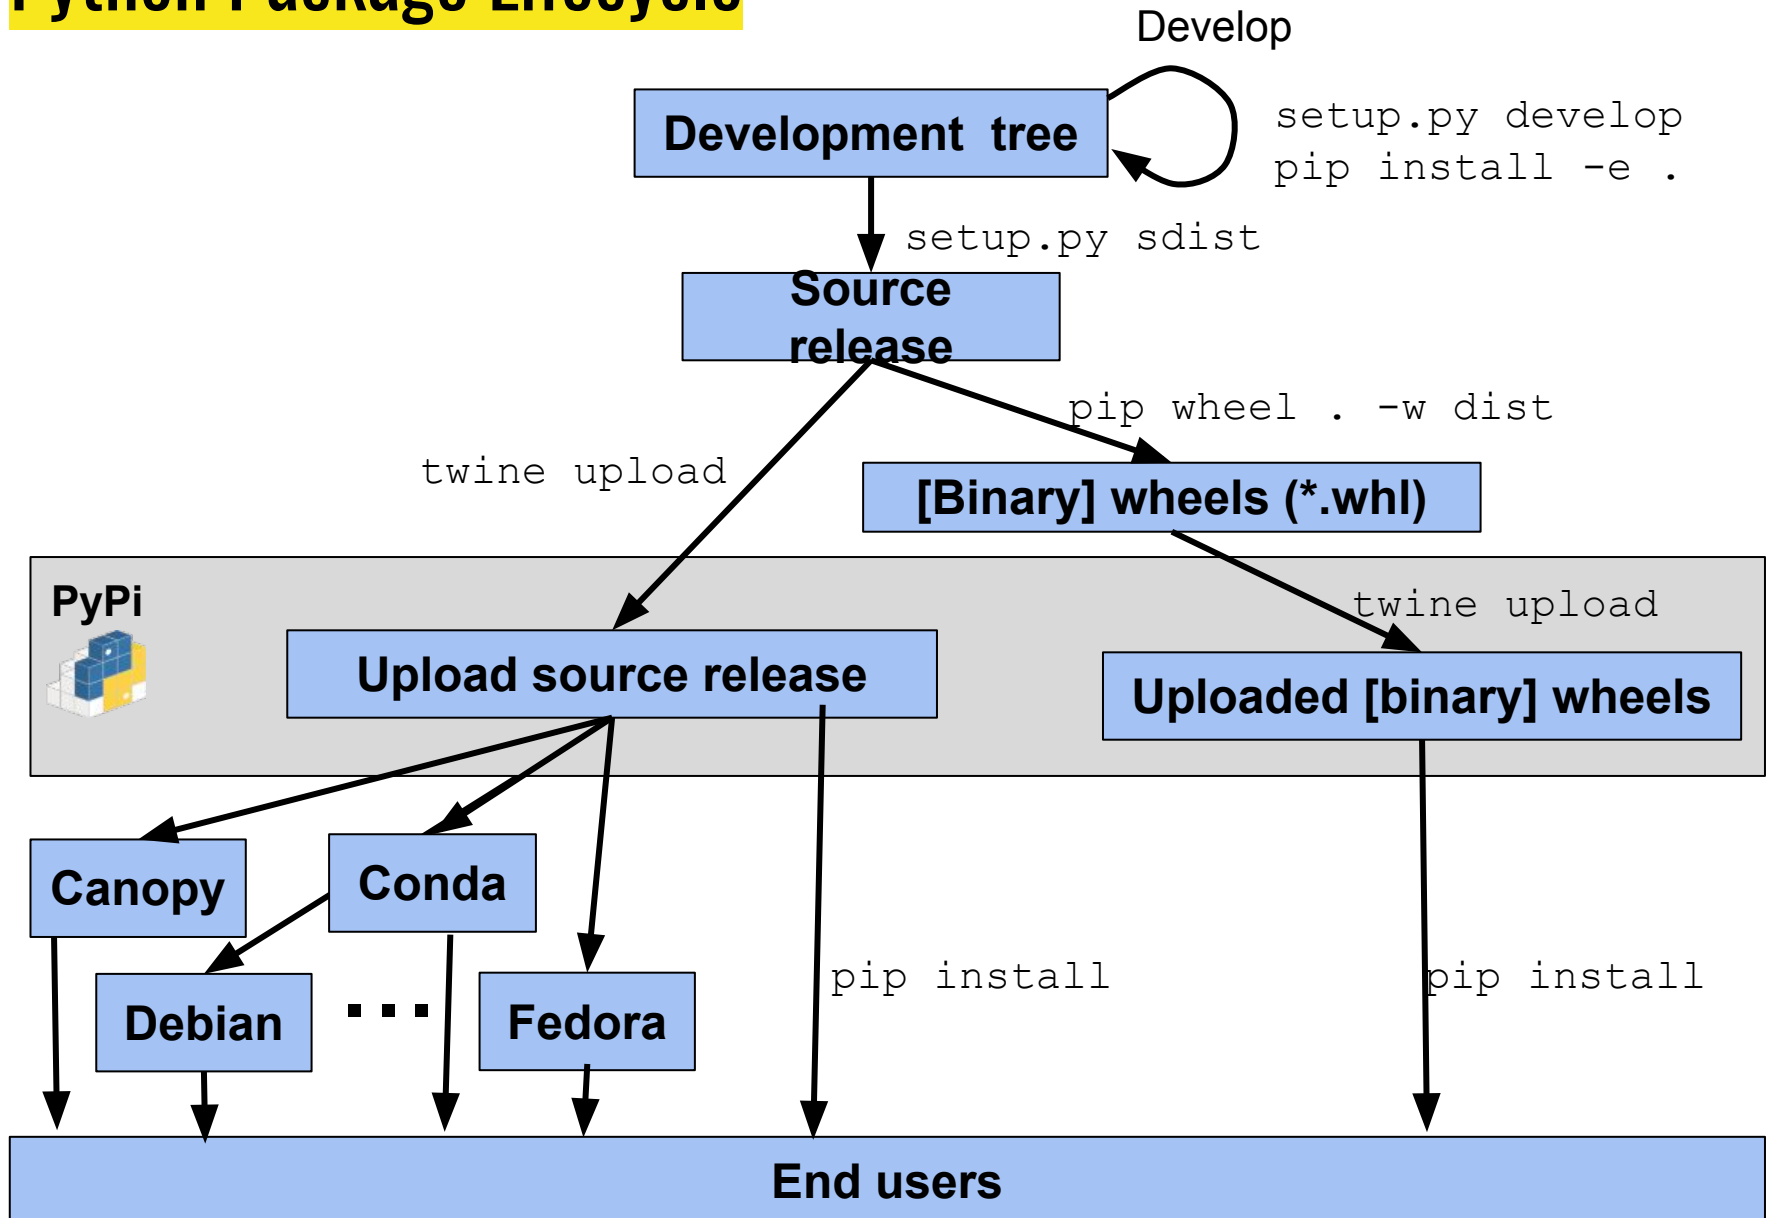
- Python Package Lifecycle
- C/C++/Fortran/Cython Binary Extensions: Problems ? Impact ?
- State of the current building tools
- Requirements
- Our solution: scikit-build
- Features
- What is next ?

The Python Packaging User Guide: Must read

A collection of tutorials and references to help you distribute and install Python packages with modern tools.

See <https://packaging.python.org/>

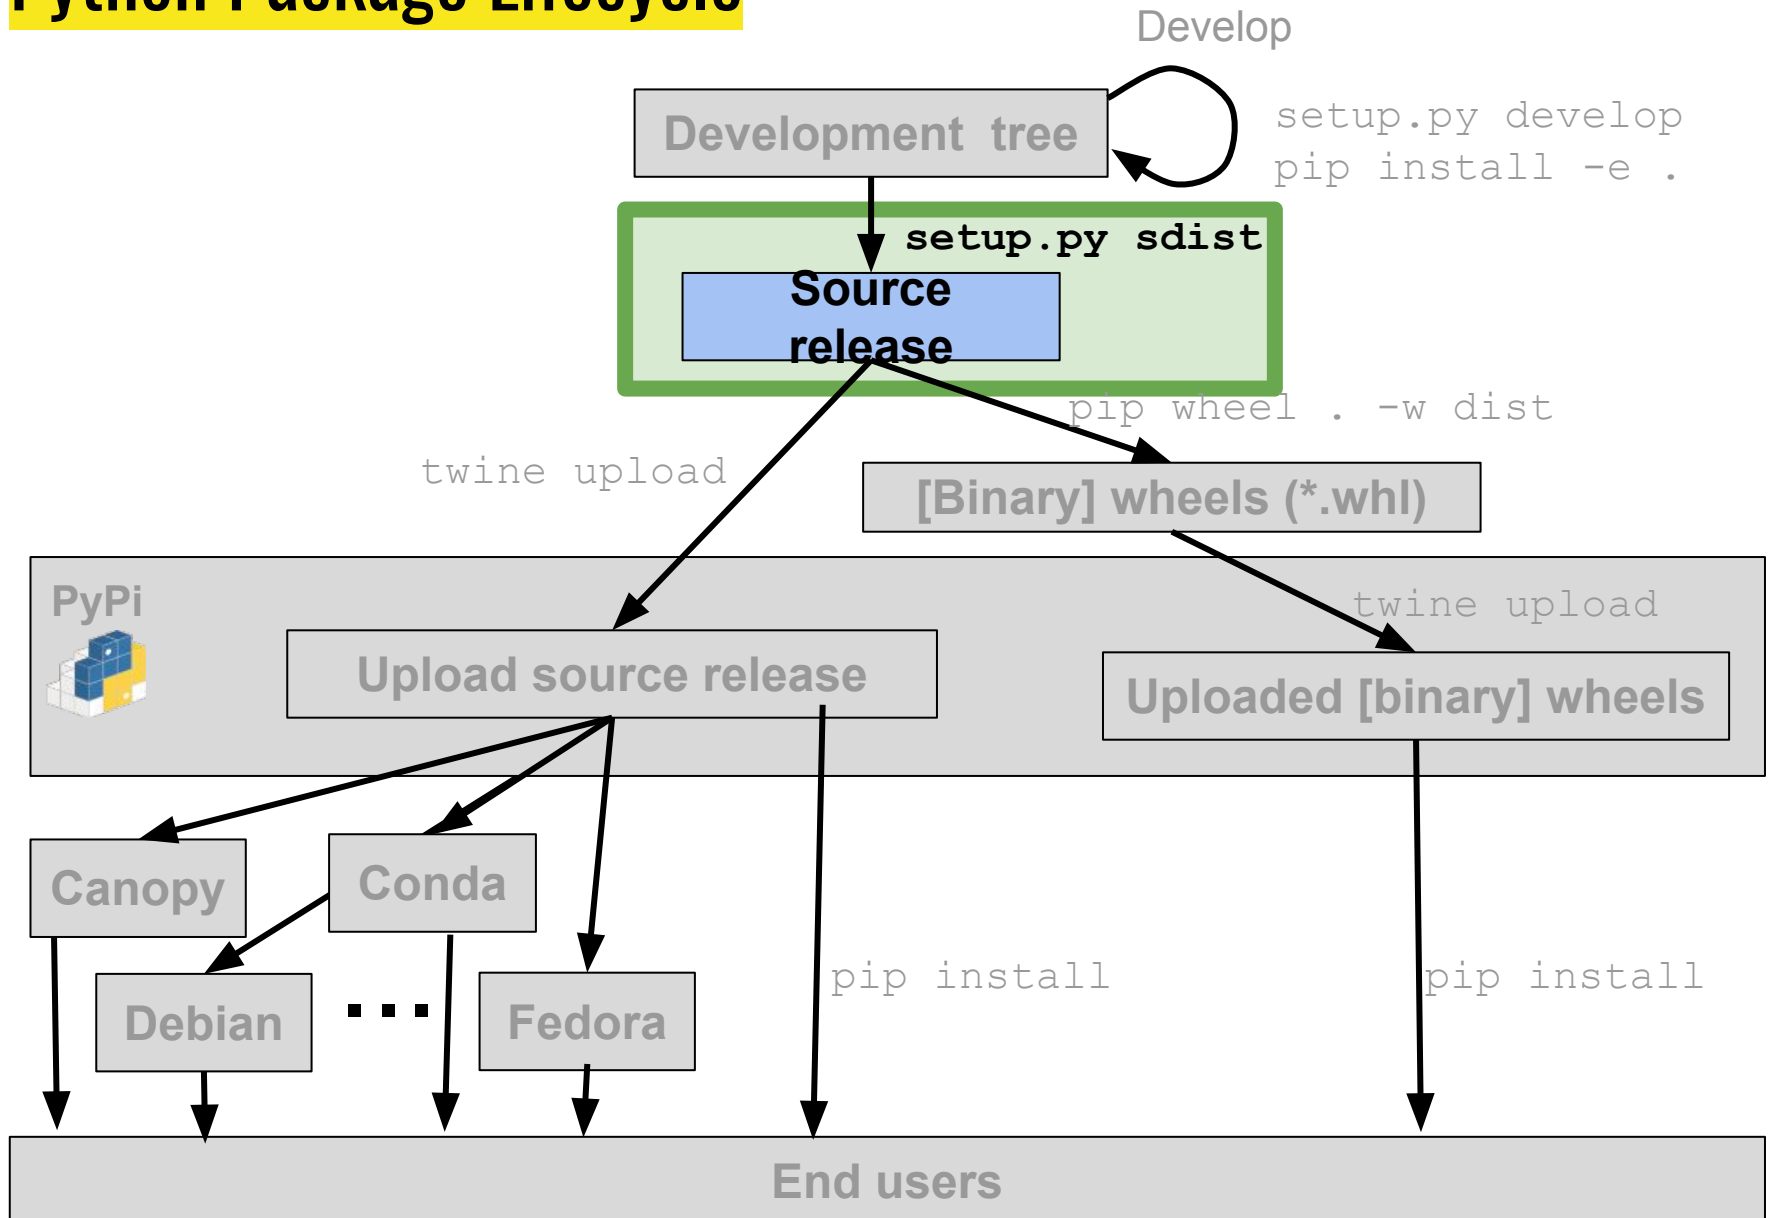
Python Package Lifecycle



Packaging Terminology 101

- **Source distribution**
 - Synonyms: **sdist**, source release
 - provides metadata + source files
 - needed for installing
 - by a tool like pip
 - or for generating a **Built Distribution**

Python Package Lifecycle

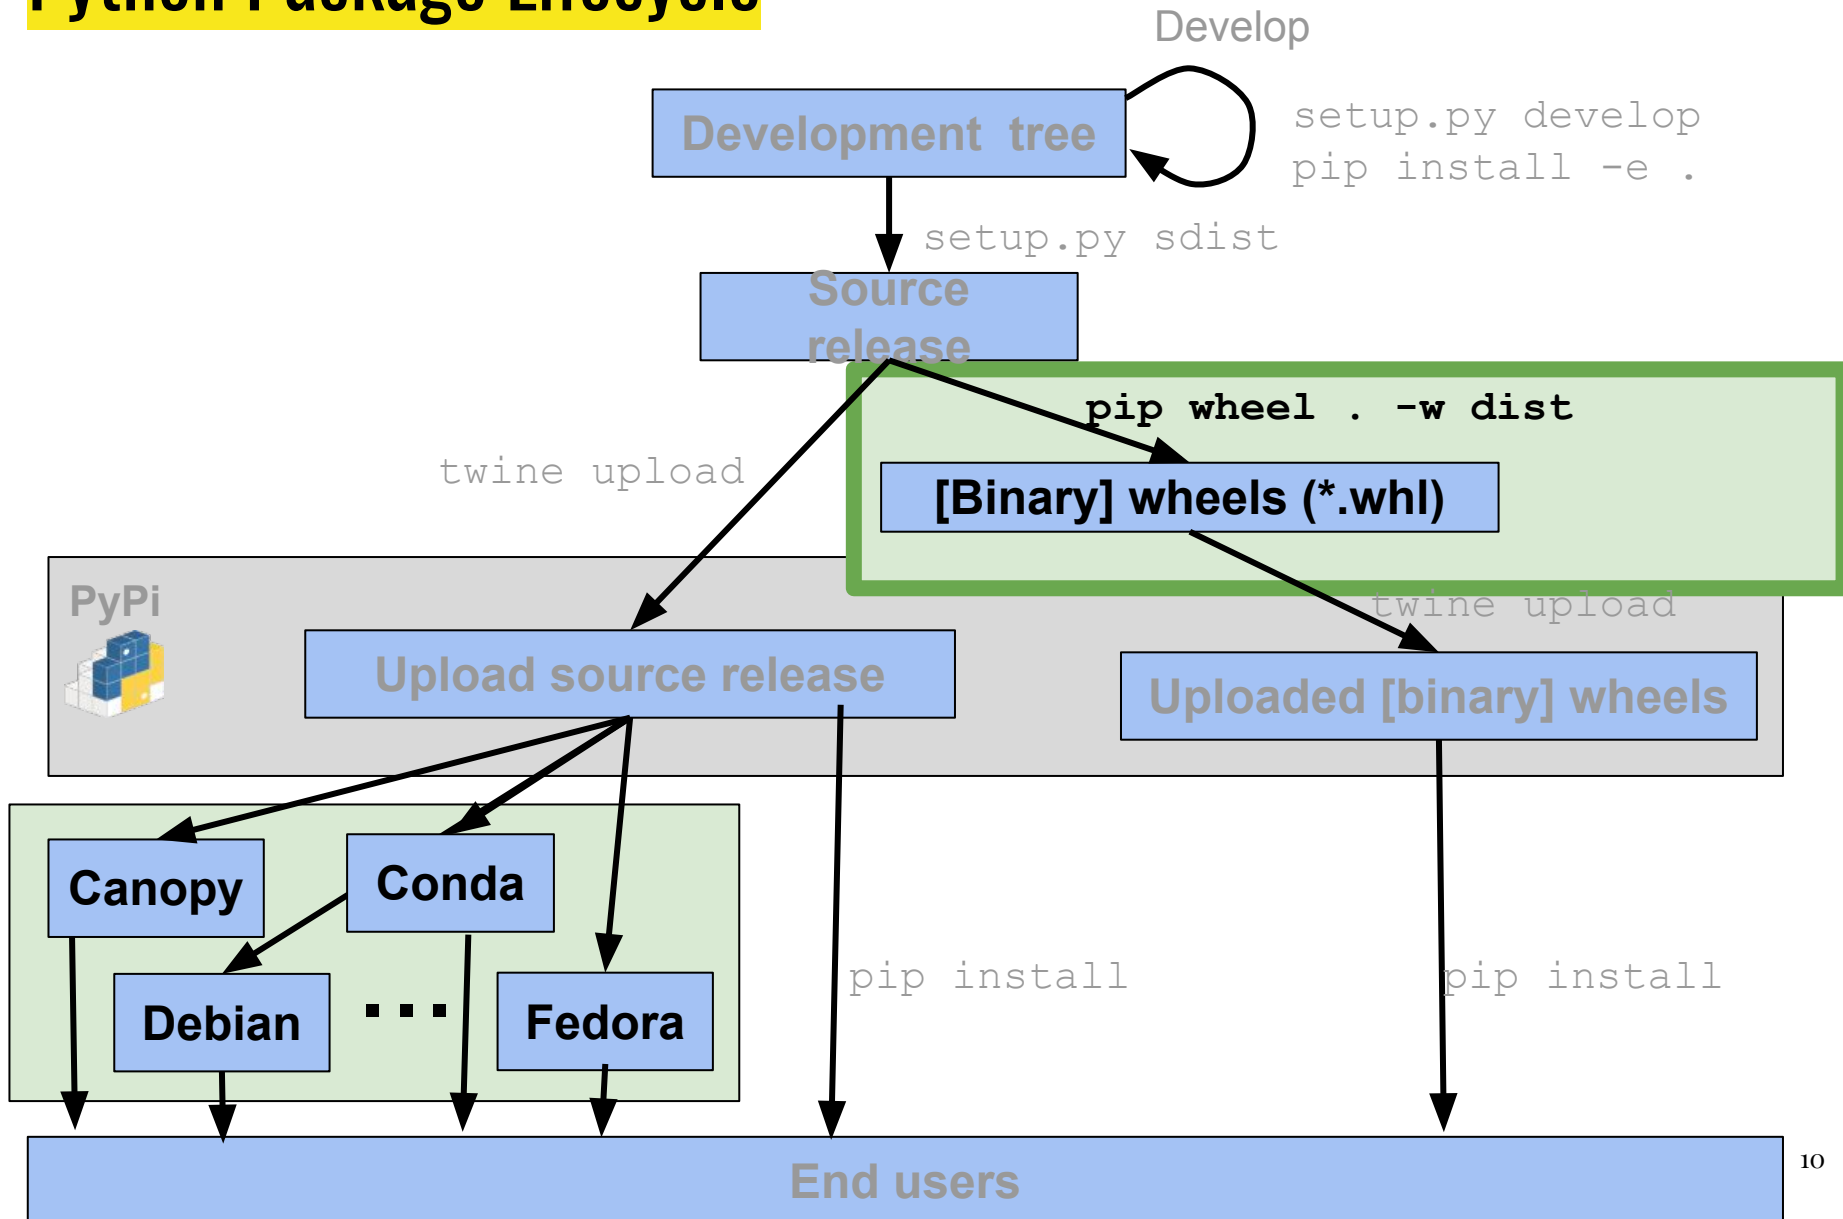


Packaging Terminology 101

- **Built Distribution**

- Synonyms: **bdist**, wheel
- provides metadata + pre-built files
- only need to be moved to the correct locations on the target system

Python Package Lifecycle



Packaging Terminology 101

- Application binary interface (ABI)
 - Interface between two binary program modules
 - Adhering to an ABI is usually the job of compiler

Packaging Terminology 101

- A python Distribution is either:
 - **pure:**
 - Not specific to a CPU architecture
 - No ABI
 - **non-pure**
 - ABI
 - Platform specific

Packaging Terminology 101

- **Binary Distribution**

- is a **Built Distribution**
- Is **non-pure**
- uses platform-specific compiled extensions

- **Non-pure**

- ABI
- Platform specific

- **Built Distribution**

- Synonyms: **bdist**, wheel
- provides metadata + pre-built files
- only need to be moved to the correct locations on the target system

Packaging Terminology 101

- **Wheel:**

- a **Built Distribution**
- a ZIP-format archive with `.whl` extension

■ `{distribution}-{version}(-{build tag})?-{python tag}-{abi tag}-{platform tag}.whl`

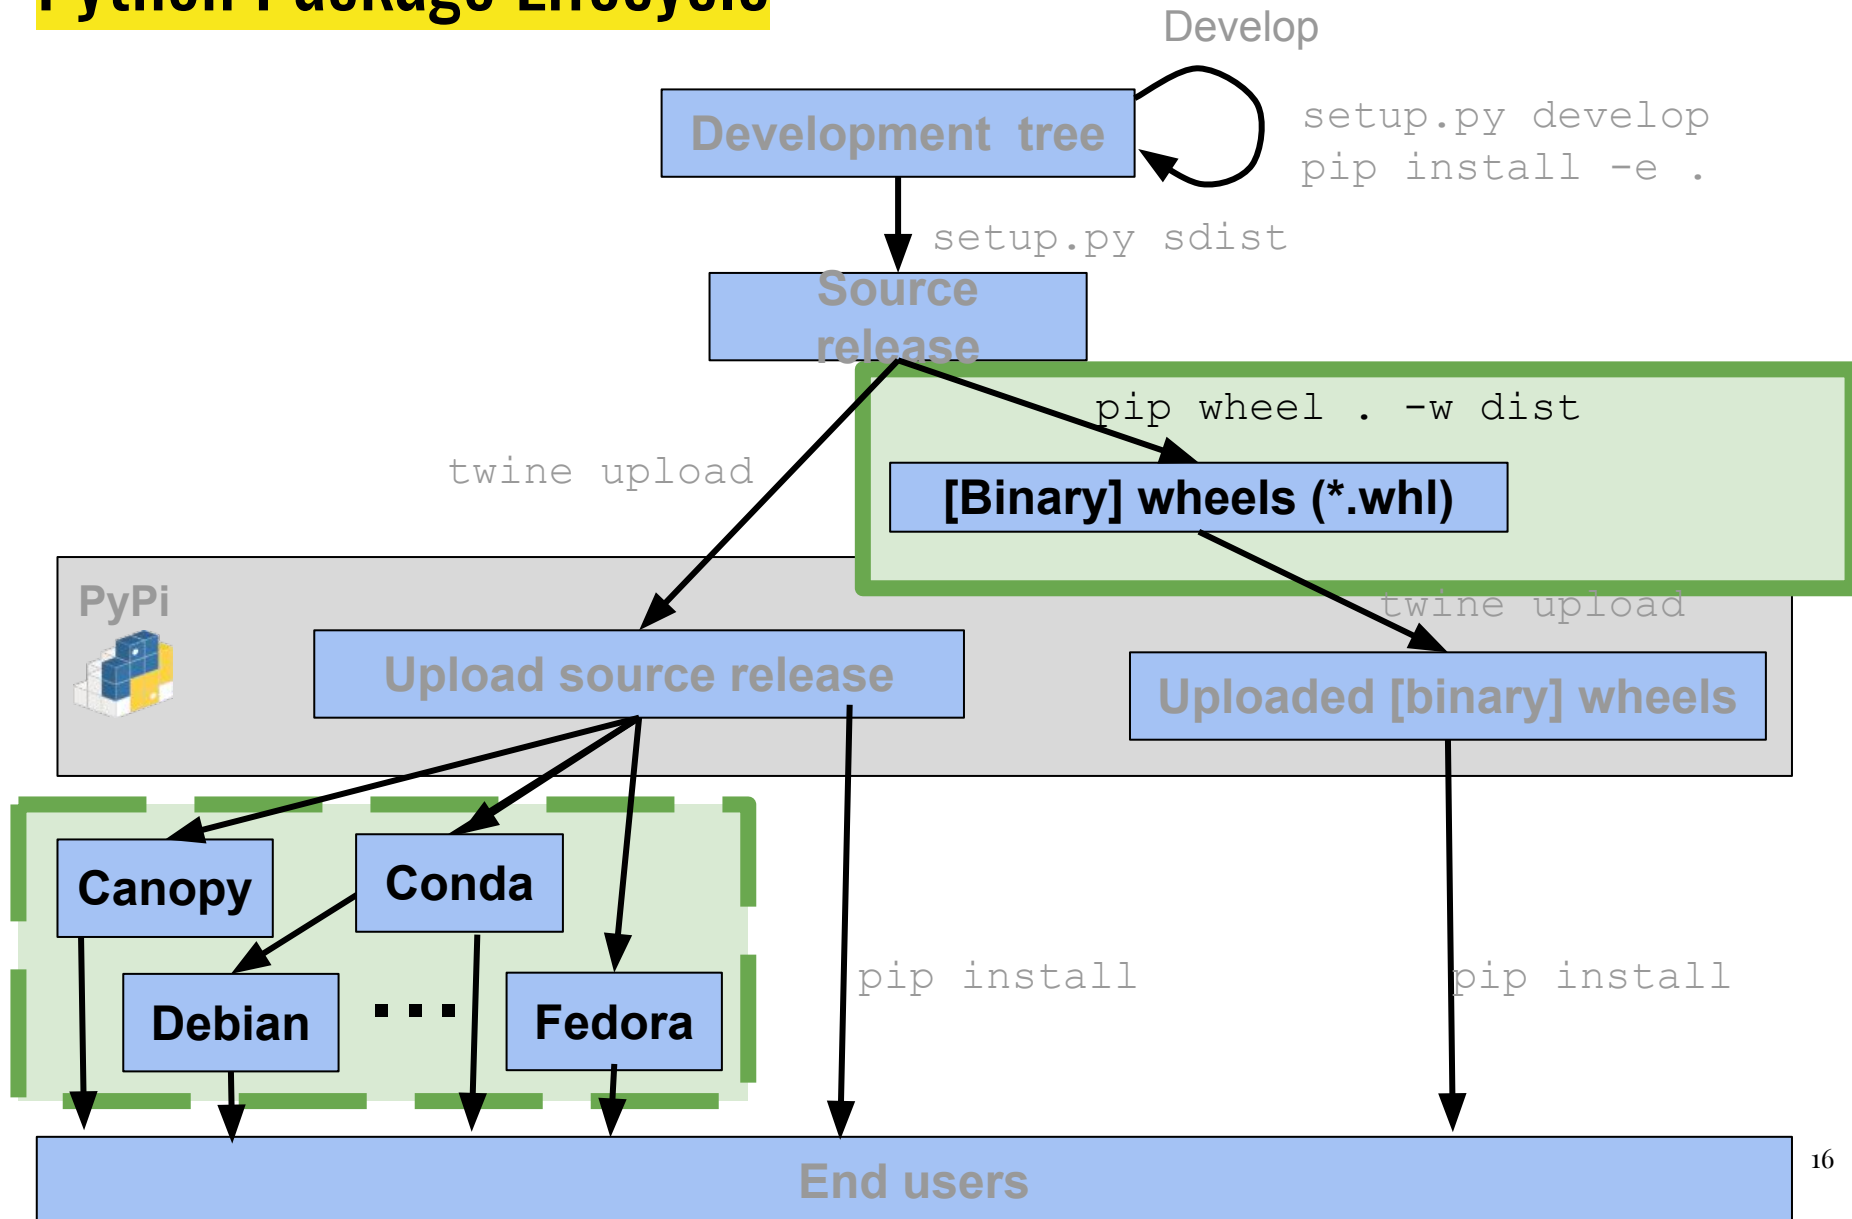
- described by [PEP 427](https://peps.python.org/pep-0427/)



Packaging Terminology 101


- Few examples of **wheel** archives:
 - **Non-pure** wheels
 - cffi-1.11.5-cp27-cp27m-manylinux1_x86_64.whl
 - cffi-1.11.5-cp27-cp27m-macosx_10_6_intel.whl
 - cffi-1.11.5-cp36-cp36m-win_amd64.whl
 - **Pure** wheels
 - docutils-0.14-py2-none-any.whl
 - docutils-0.14-py3-none-any.whl
 - certifi-2018.4.16-py2.py3-none-any.whl

Python Package Lifecycle



The Python Packaging User Guide: Caveat

“Packaging of binary extensions” needs some TLC

 PyPA » Python Packaging User Guide » Guides »

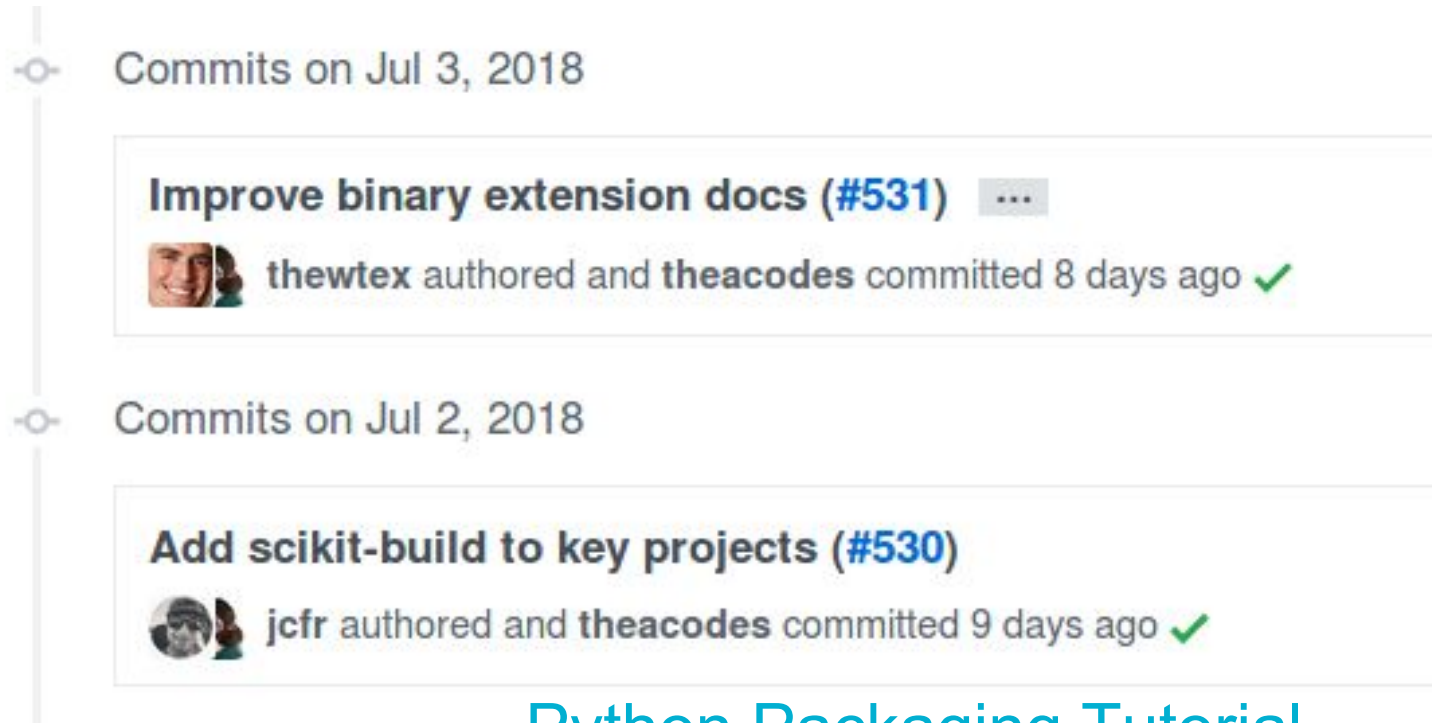
Packaging binary extensions



Page Status:	Incomplete
Last Reviewed:	2013-12-08

One of the features of the CPython reference interpreter is that, in addition to allowing the execution of Python code, it also exposes a rich C API for use by other software. One of the most common uses of this C API is to create importable C extensions that allow things which aren't always easy to achieve in pure Python code.

The Python Packaging User Guide: Good news



[Python Packaging Tutorial](#)

Thanks @theacodes for reviewing and integrating !

The SciPy Python Packaging Tutorial

The Sheer Joy of Packaging!

Scipy 2018 Tutorial

Packaging

Packaging from start to finish for both PyPI and conda

Topics

- Tutorial Schedule
- Overview
- Making a Python Package
- Building and Uploading to PyPI
- Binaries and Dependencies
- Conda Packages

<https://python-packaging-tutorial.github.io/python-packaging-tutorial/>

C/C++/Fortran/Cython Binary Extensions

Question:

Do they really matter for science ?

C/C++/Fortran/Cython Binary Extensions

Question:

Do they really matter for science ?

Answer:

Yes, they allow to:

- speed up computationally intensive code
- create wrapper module
- access lower level features of OS, hardware, CPython, ...
- leverage multi-core, many-core, GPU, ...

C/C++/Fortran/Cython Binary Extensions

Question:

Do they really matter for science ?

Answer:

Yes, they allow to:

- ...
- leverage existing scientific code
- careful and precise memory management
- avoid parallel computing issues related to CPython's [Global Interpreter Lock \(GIL\)](#)

C/C++/Fortran/Cython Binary Extensions

Question:

Sure, but which packages include them ?

C/C++/Fortran/Cython Binary Extensions

Question:

Sure, but which packages include them ?

Answer:

Just to name a few ... cffi, cython, lxml, matplotlib, msgpack-python, numpy, pandas, pillow, pyaml, pyne, pyzmq, scikit-learn, **scipy**, symengine, twisted, mayavi

What are the problems ?

- Specifying the correct compiler and link flags is error-prone
- Finding build-time dependencies is difficult



What are the problems ?

- Leveraging build tool like make or ninja is not possible
- Building multiple extensions in parallel is not an option without some crazy [hack](#)
- Using an IDE like Visual Studio, QtCreator, ... to develop the C or C++ code is not possible



What are the problems ?

- Setting up a project with compiled modules is tedious
- “Cross-platform development and distribution of extension modules is a **complex topic** [...]”

Source: [Python Packaging User Guide / Packaging binary extensions](#)

- Cross-compiling is horrendous



Who does it impact ?

- You writing code
- You maintaining a project
- You installing a “package”

There is now “scikit-build”

A **reliable** and **easy** way to package your **compiled** Python projects.



State of the current tools

- handle common needs
- but leave a lot to be desired
- not specialized for compiling

Requirements (1 / 3)

- reliably building native codes
- first class support for system introspection
- creating extension modules for Python scripts

Requirements (2 / 3)

- dynamically and statically linking extension modules for use in embedded applications
- compiling, packaging and publishing
- all integrated with the Python ecosystem

Requirements (3 / 3)

- first-class cross platform support
- cross-compilation capabilities

Observation

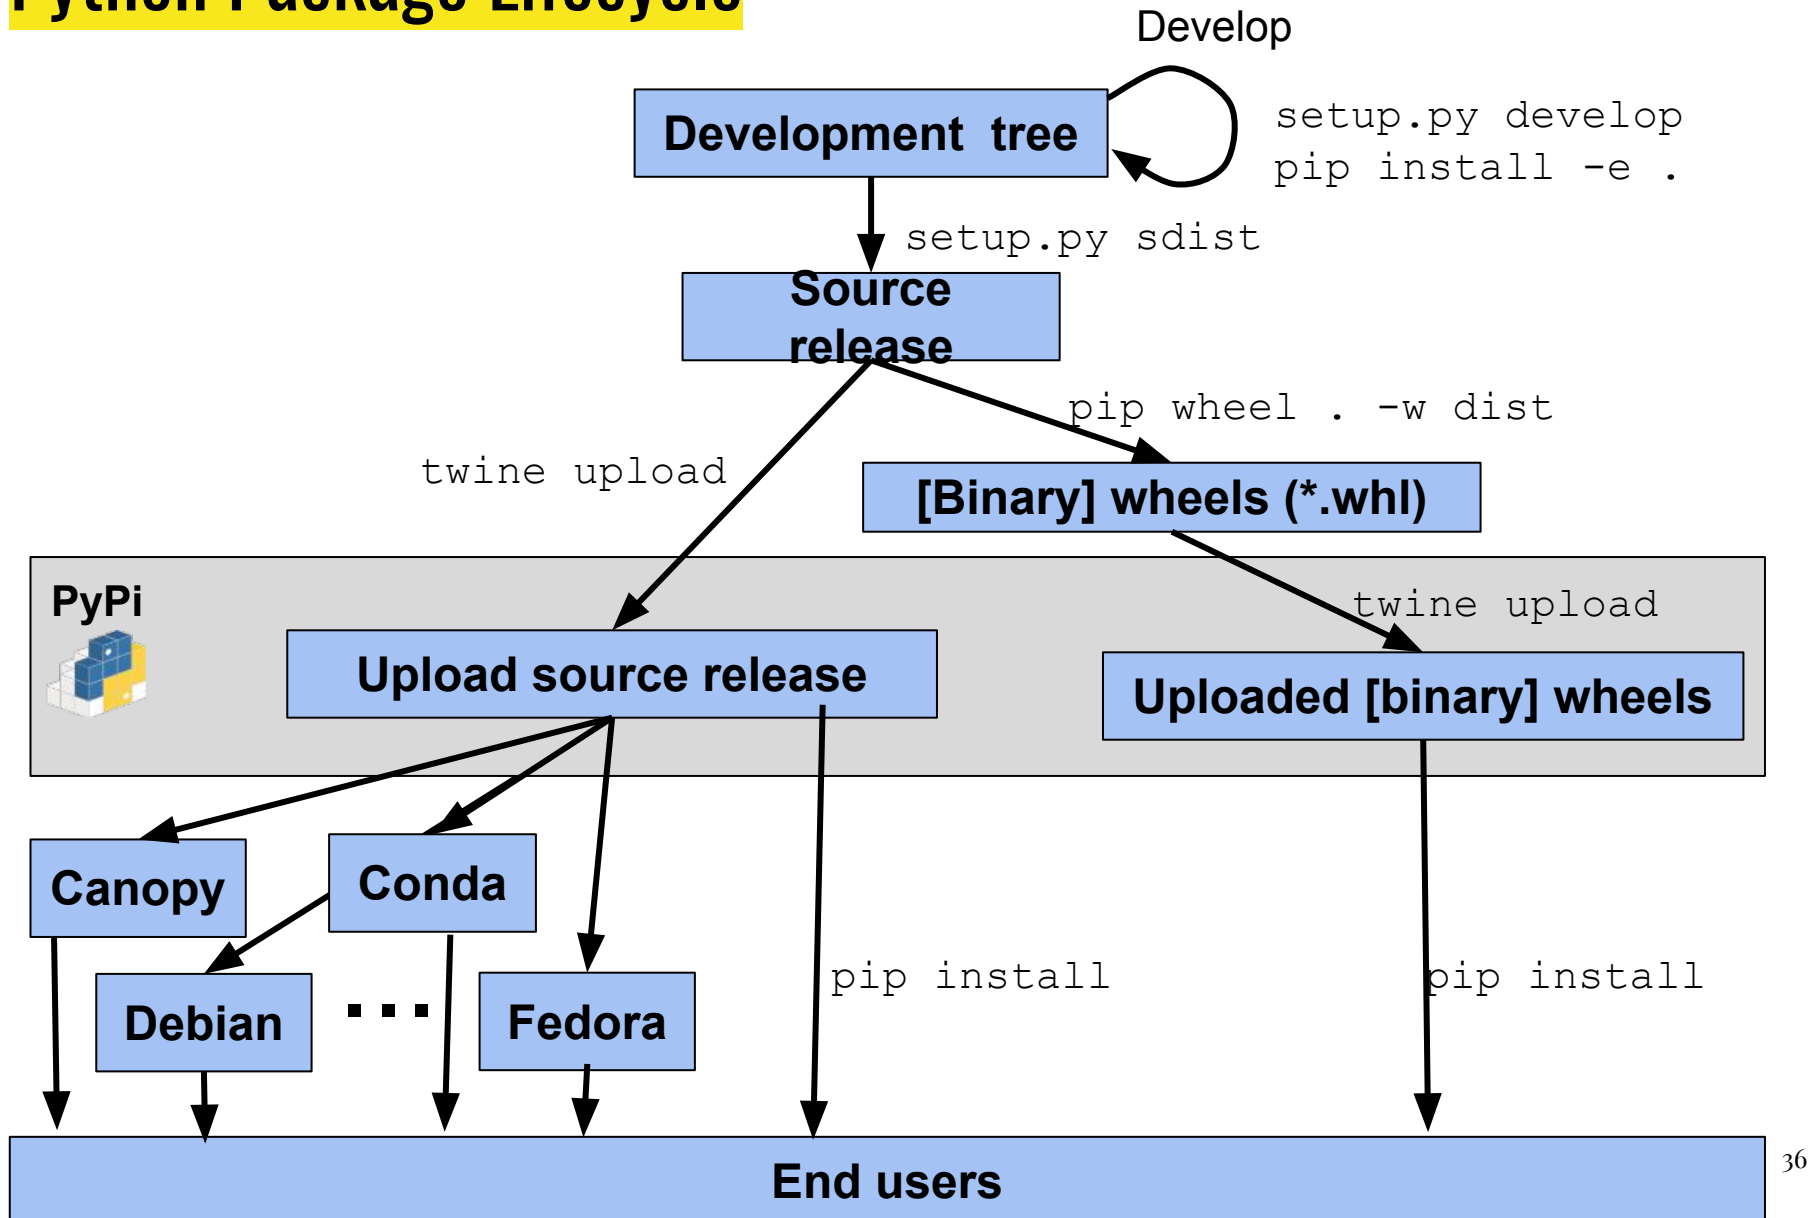
There is a clear need for a **cross-platform** building and packaging solution that **supports** projects using **compiled** modules and their **users**.

Our solution: scikit-build

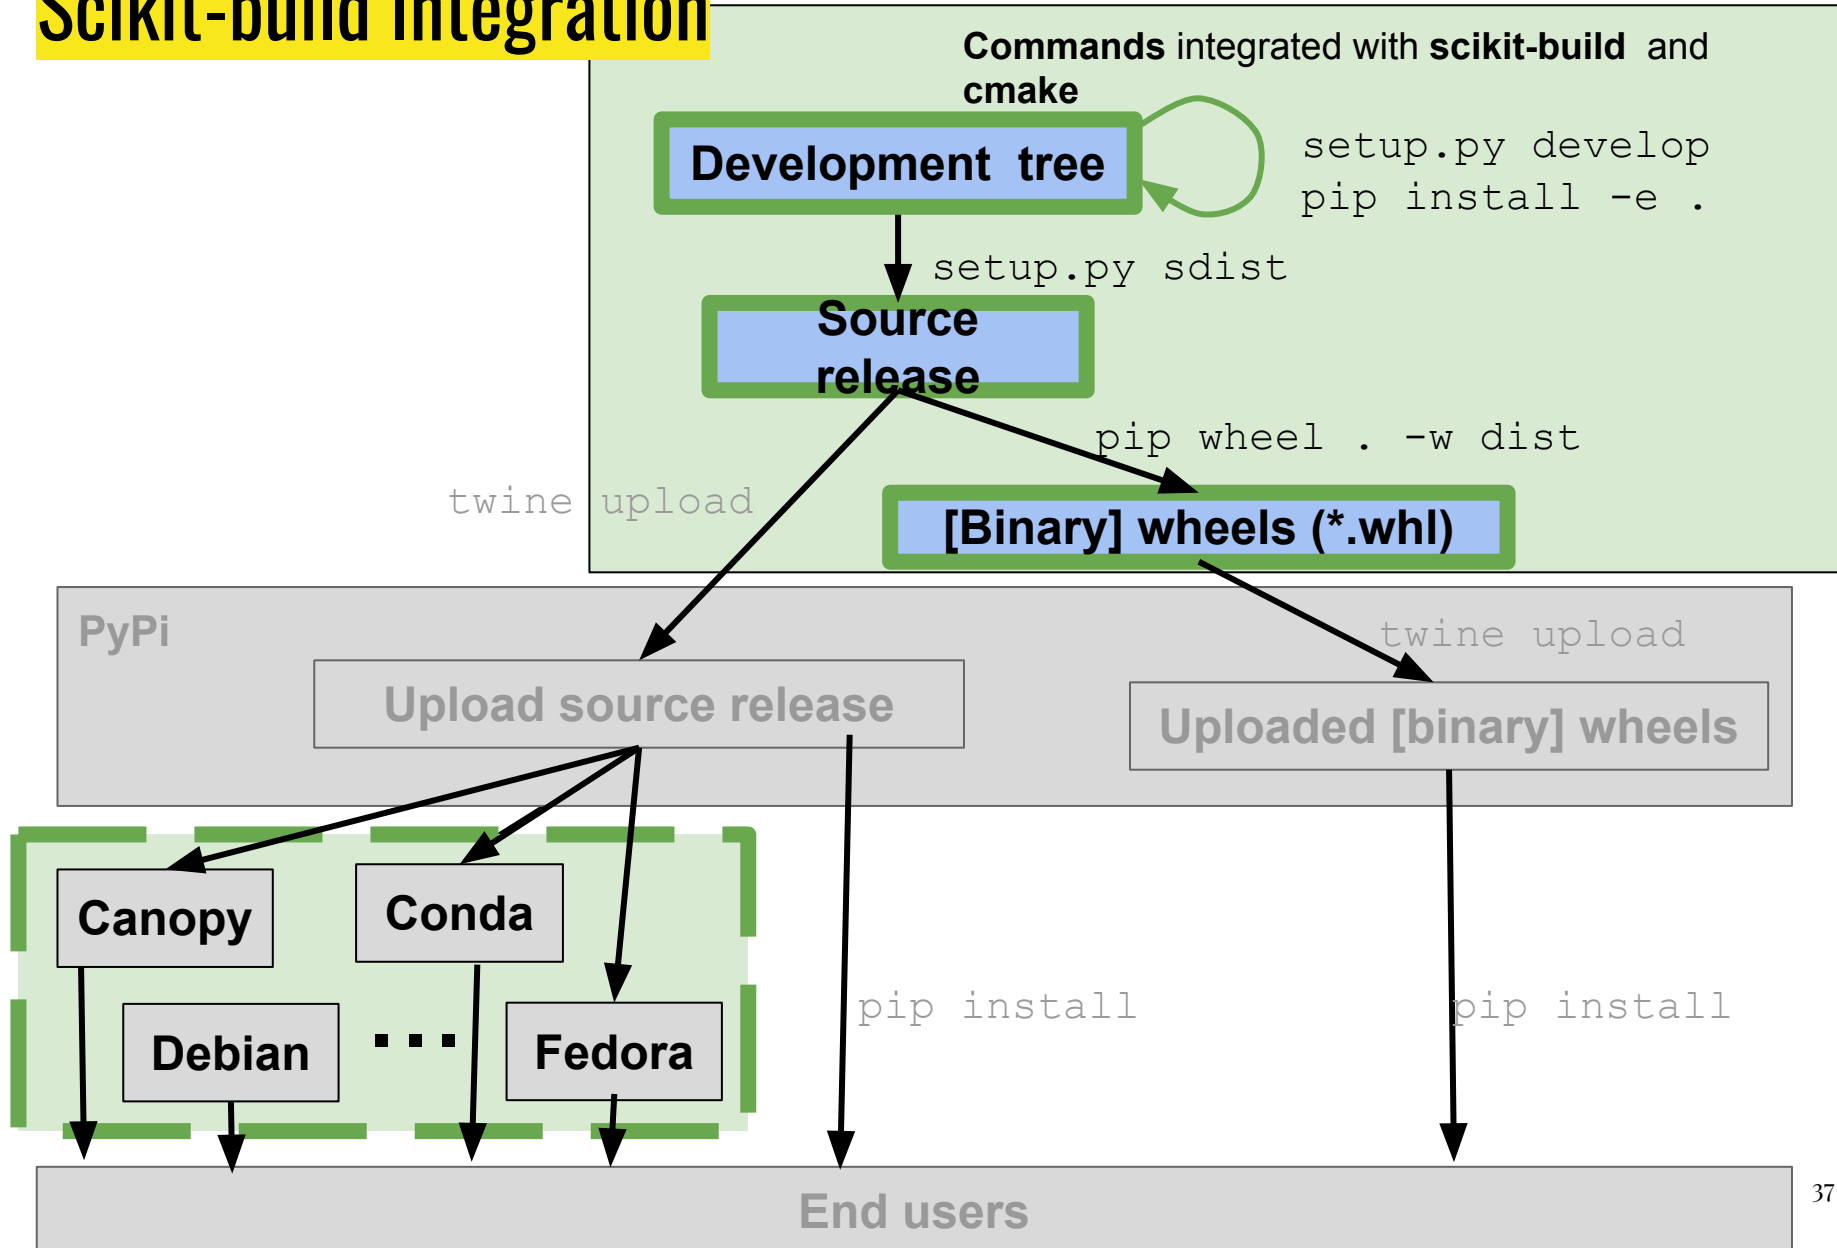
Packaging solution that **bridges the gap** between

- [cmake](#),
- **your** C/C++/Fortran and Cython **projects**,
- and the **Python ecosystem** ([pip](#) and [setuptools](#))

Python Package Lifecycle




Scikit-build Integration



Benefits of using CMake

- open-source, well-maintained and supported
- cross-platform builds with support for cross-compilation
- system introspection capabilities
- support for native build systems
- python extension module support

Build Tool Trends



Source:

https://trends.google.com/trends/explore?cat=5&date=all&q=%2Fm%2F0cxh7f,%2Fg%2F1thtpq_p,%2Fm%2F01fg0l,%2Fg%2F11bzyq50jp,%2Fm%2F04d104

scikit-build is ...

- **not** (yet) reinventing the .whl
- **not** a new paradigm:
 - It integrates with ***existing*** tools
- **not** using `setuptools.Extension` internally but can be used along side

scikit-build is ...

- **not** a replacement for canopy, conda or pypi
 - does **not** manage packages
 - only makes building them easier
- **not** magic
 - you're still compiling extension modules
 - only with much less guess-work

scikit-build is ...

- a drop-in replacement for setuptools

```
--- a/setup.py
```

```
+++ b/setup.py
```

```
-from setuptools import setup
```

```
+from skbuild import setup
```

scikit-build only needs ...

- a `CMakeLists.txt` describing your extension

Hello Example: Layout

```
$ cd hello && find .  
.  
./CMakeLists.txt  
./hello  
./hello/_hello.cxx  
./hello/__init__.py  
./setup.py  
./pyproject.toml
```

Hello Example: setup.py

```
from skbuild import setup

setup(
    name="hello",
    version="1.2.3",
    packages=['hello'],
    # description, author, license, ...
)
```

Hello Example: CMakeLists.txt

```
cmake_minimum_required(VERSION 3.12)
project(hello)
find_package(PythonExtensions REQUIRED)
add_library(_hello MODULE hello/_hello.cxx)
python_extension_module(_hello)
install(TARGETS _hello LIBRARY DESTINATION hello)
```

Hello Example: pyproject.toml

```
[build-system]
requires = ["setuptools", "wheel", "scikit-build",
"cmake", "ninja"]
```

Hello Example: Building the wheel

```
$ cd hello
```

```
$ pip wheel . -w dist
```


Scikit-build Features (1 / 3)

- Support for Python 2.7, 3.4 and above
- Compiling environment
 - By default, lookup compilers matching python version
 - Provide recommendations if compilers not installed
 - Understand CC, CXX, CFLAGS and CXXFLAGS env. variables
- Support ninja build tool on Linux, macOS and Windows
 - No need to toy around with vcvarsall.bat et al

Scikit-build Features (2 / 3)

- Support for developer mode: `pip install -e .`
- [cmake](#) and [ninja](#) pip installable on all platforms
- Support generation of Source Distribution for git repository (only if `MANIFEST.in` is not found)
- Fast incremental rebuild by skipping re-configuration if relevant.

Scikit-build Features (2 / 3)

- Support for **setup usual keywords**: packages, package_dir, include_package_data, package_data, exclude_package_data, py_modules, data_files, scripts
- Support for additional setup keywords to **configure CMake**: cmake_args, cmake_install_dir, cmake_source_dir, cmake_with_sdist, cmake_languages

scikit-build: Command line options

- Useful for passing options
 - to CMake
 - to build tool (make, ninja, msbuild, ...)

```
$ python setup.py \  
    bdist_wheel -- <cmake_options> -- <buildtool_options>
```

Scikit-build also provides ...

[PythonExtensions](#) CMake module to build python module or standalone python executable.

[Cython](#) CMake module to generate source code.


[NumPy](#) CMake module to lookup numpy/arrayobject.h, conv-template and from-template.

[F2PY](#) CMake module to find f2py executable facilitating creating/building of Python extension calling Fortran 77/90/95 external subroutines.

Scikit-build: Software Process, CI and Documentation

- ~200 tests running on 3 platforms for python 2.7, 3.4 to 3.7
- <https://scikit-build.readthedocs.io>
- [CONTRIBUTING](#), [Release Notes](#), MIT license, [mailing list](#)

Build Status

	Linux	MacOSX	Windows
PyPI	 circleci passing	 build passing	 build passing
Conda	 circleci passing	 build passing	 build passing

Overall Health

requirements	up-to-date	codecov	90%	health	99%
--------------	------------	---------	-----	--------	-----











What is next ?

- Scipy2018 sprints (Saturday):
 - finalize conda package associated with latest release 0.7.1
 - improve documentation and add more example projects
 - provide guidance to integrate scikit-build
- Long term
 - work with numpy and scipy folks
 - help with [proposal](#) for packaging native libraries into Python wheels
 - implement as [PEP 517](#) build backend interface, configured with pyproject.toml and CMakeLists.txt ([issue #124](#))
 - leverage [distlib](#) (library of packaging functionality)

Thank you !

Table generated using [mgechev/github-contributors-list](#) npm package


- scikit-build contributors:

				
willingc	seanlis	neok-m4700	henryiii	blowekamp
				
scopatz	msmolens	thewtex	msarahan	opadron

and also xoviat

- Scikit-build issue reporters: benjaminjack, isuruf, henryborchers, jonwoodring, mivade, reiver-dev, seanlis
- PyPA, PEPs contributors, the wider community

[Attribution 4.0 International](#)

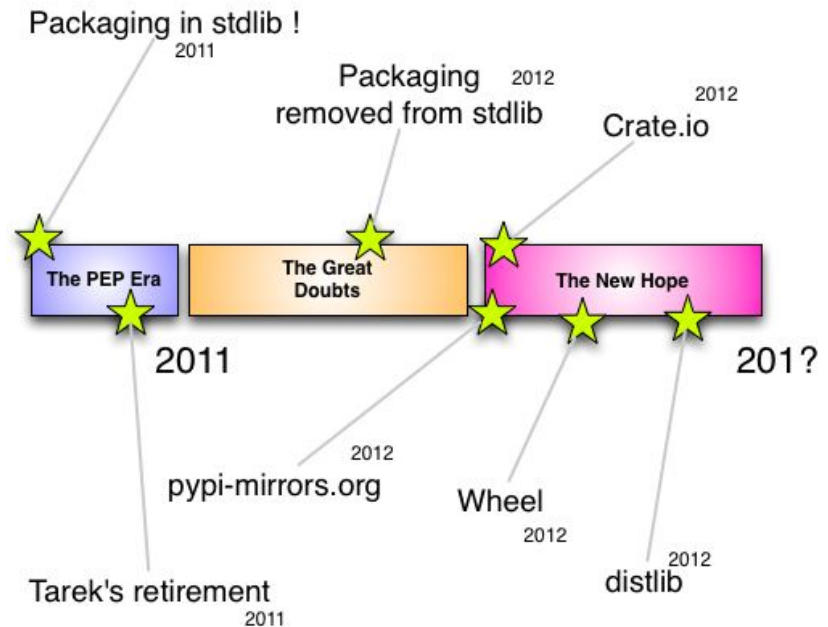
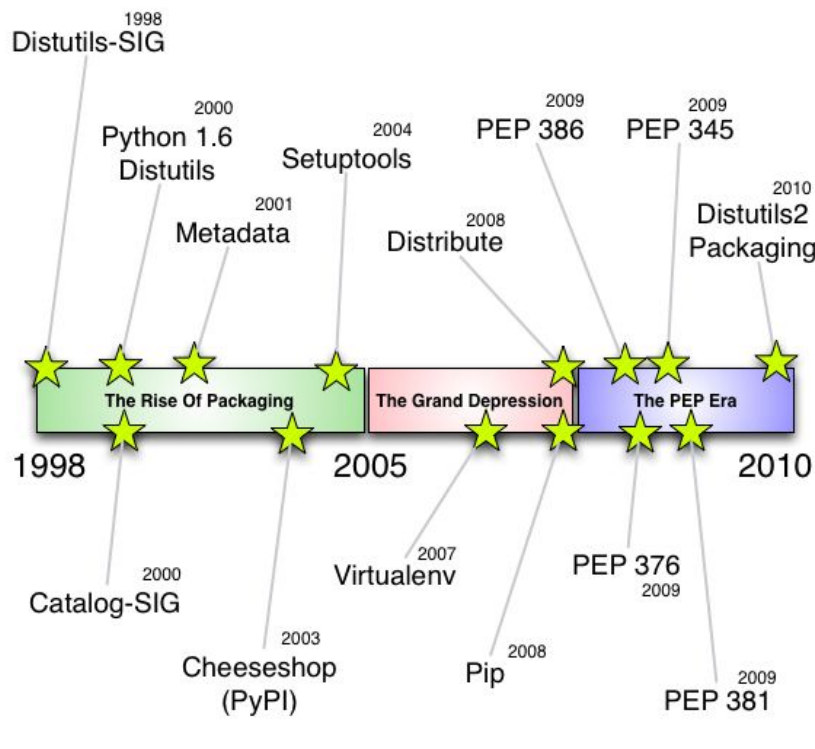


scikit-build

Slides: bit.ly/scikit-build-talk

Source code: github.com/scikit-build/scikit-build

Python Packaging History: A long standing community effort



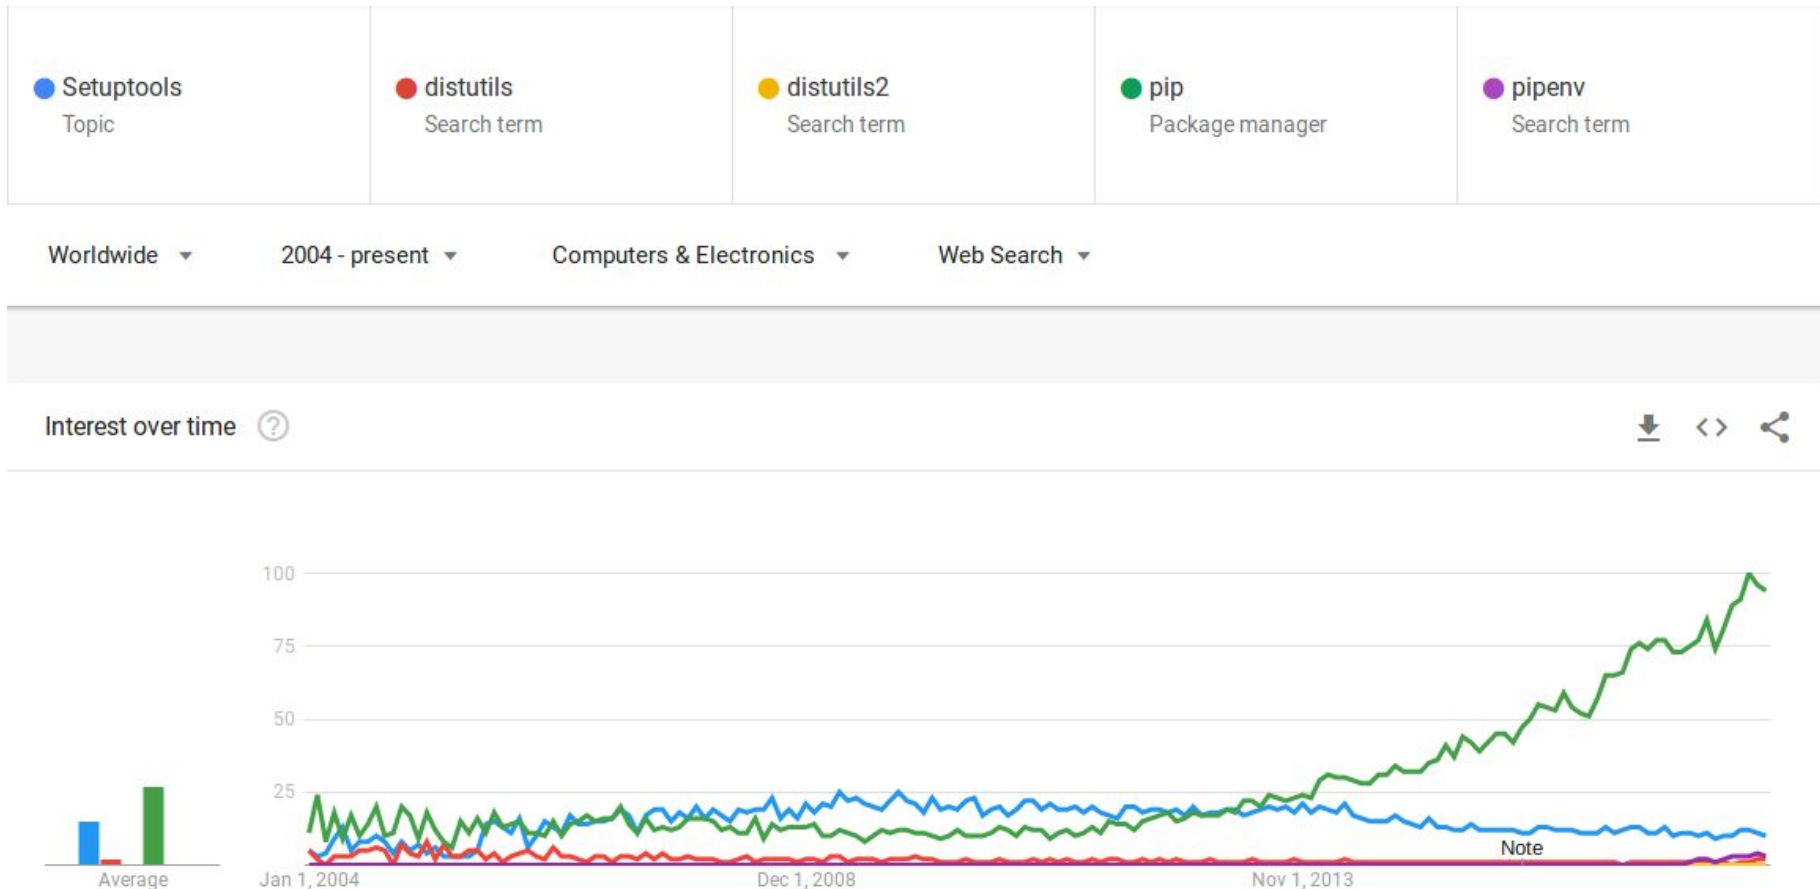
Source: <https://ziade.org/2012/11/17/chronology-of-packaging/>

Followed by outstanding standardization efforts

- 2013:
 - [PEP 425](#) -- Compatibility Tags for Built Distributions
 - [PEP 427](#) -- The Wheel Binary Package Format 1.0
- 2014:
 - [PEP 440](#) -- Version Identification and Dependency Specification
- 2016:
 - [PEP 513](#) -- A Platform Tag for Portable Linux Built Distributions
 - [PEP 518](#) -- Specifying Minimum Build System Requirements for Python Projects
- 2017:
 - [PEP 517](#) -- A build-system independent format for source trees

and many more at <https://www.python.org/dev/peps/>

Python Packaging History: The trend



Source: <https://trends.google.com/trends/explore?cat=5&date=all&q=%2Fm%2F07kg5hp.distutils.distutils2.%2Fm%2F0hzm844.pipenv>

- [Pip](#): \$778k, 14 person-years, The python package manager
- [Setuptools](#): \$48k, 1 person-year, enhancements to the Python distutils
- [Scikit-build](#): \$65k, 1 person-year » Our tool, simple and to the point. It creates a bridge.
- [Python](#): \$15M, 284 person-years » CPython interpreter
- [CMake](#): \$19M, 346 person-years » Specialized in compiling and building