




*Fast Data Structures for Data Sciences*

# What is xtensor?

A modern C++ template Library for fast N-D tensor algebra. TLDR: NumPy for C++



Language bindings, to operate *in-place* on the data structures of the main languages of data science.



Solid foundation meant for *speed* and ease of use



But Also:

FFT, Linear Algebra bindings



Utilities to read & write various file formats from and to xtensor (numpy/npz, HDF5, image files, audio files)



A C++ data frame in the making



An much more! ROS bindings, a cookiecutter project for authoring Python or Julia extensions...

# From NumPy to xtensor

Cheat sheet available at <http://xtensor.readthedocs.io/>

Longer examples from NumPy translated in <https://github.com/wolfv/numpy-benchmarks>

## Containers

### Python 3 - NumPy

```
np.array([[3, 4], [5, 6]])
```

```
arr.reshape([3, 4])
```

```
arr.astype(np.float64)
```

### C++ 14 - xtensor

```
xt::xarray<double>({{3, 4}, {5, 6}})
```

```
xt::xtensor<double, 2>({{3, 4}, {5, 6}})
```

```
arr.reshape({3, 4})
```

```
xt::cast<double>(arr)
```

## Initializers

```
np.linspace(1.0, 10.0, 100)
```

```
np.logspace(2.0, 3.0, 4)
```

```
np.arange(3, 7)
```

```
np.eye(4)
```

```
np.zeros([3, 4])
```

```
np.ones([3, 4])
```

```
np.empty([3, 4])
```

```
np.meshgrid(x0, x1, x2, indexing='ij')
```

```
xt::linspace<double>(1.0, 10.0, 100)
```

```
xt::logspace<double>(2.0, 3.0, 4)
```

```
xt::arange(3, 7)
```

```
xt::eye(4)
```

```
xt::zeros<double>({3, 4})
```

```
xt::ones<double>({3, 4})
```

```
xt::empty<double>({3, 4})
```

```
xt::meshgrid(x0, x1, x2)
```

## Broadcasting

### Python 3 - NumPy

```
a[:, np.newaxis]  
a[:5, 1:]  
a[5:1:-1, :]  
a[..., 3]
```

```
np.broadcast(a, [4, 5, 7])  
np.vectorize(f)  
a[a > 5]  
a[[0, 1], [0, 0]]
```

### C++ 14 - xtensor

```
xt::view(a, xt::all(), xt::newaxis())  
xt::view(a, xt::range(_, 5), xt::range(1, _))  
xt::view(a, xt::range(5, 1, -1), xt::all())  
xt::strided_view(a, {xt::ellipsis, 3})  
  
xt::broadcast(a, {4, 5, 7})  
xt::vectorize(f)  
xt::filter(a, a > 5)  
xt::index_view(a, {{0, 0}, {1, 0}})
```

## Random

```
np.random.seed(0)  
np.random.randn(10, 10)  
np.random.randint(10, 10)  
np.random.rand(3, 4)  
np.random.choice(arr, 5)
```

```
xt::random::seed(0)  
xt::random::randn<double>({10, 10})  
xt::random::randint<int>({10, 10})  
xt::random::rand<double>({3, 4})  
xt::random::choice(arr, 5)
```

## Broadcasting

### Python 3 - NumPy

```
a[:, np.newaxis]  
a[:5, 1:]  
a[5:1:-1, :]  
a[..., 3]
```

```
np.broadcast(a, [4, 5, 7])  
np.vectorize(f)  
a[a > 5]  
a[[0, 1], [0, 0]]
```

### C++ 14 - xtensor

```
xt::view(a, xt::all(), xt::newaxis())  
xt::view(a, xt::range(_, 5), xt::range(1, _))  
xt::view(a, xt::range(5, 1, -1), xt::all())  
xt::strided_view(a, {xt::ellipsis, 3})  
  
xt::broadcast(a, {4, 5, 7})  
xt::vectorize(f)  
xt::filter(a, a > 5)  
xt::index_view(a, {{0, 0}, {1, 0}})
```

## Random

```
np.random.seed(0)  
np.random.randn(10, 10)  
np.random.randint(10, 10)  
np.random.rand(3, 4)  
np.random.choice(arr, 5)
```

```
xt::random::seed(0)  
xt::random::randn<double>({10, 10})  
xt::random::randint<int>({10, 10})  
xt::random::rand<double>({3, 4})  
xt::random::choice(arr, 5)
```

## Stack, concatenate, split

### Python 3 - NumPy

```
np.stack([a, b, c], axis=1)
np.concatenate([a, b, c], axis=1)
np.squeeze(a)
np.expand_dims(a, 1)
np.atleast_3d(a)
np.split(a, 4, axis=0)
```

### C++ 14 - xtensor

```
xt::stack(xtuple(a, b, c), 1)
xt::concatenate(xtuple(a, b, c), 1)
xt::squeeze(a)
xt::expand_dims(a, 1)
xt::atleast_3d(a)
xt::split(a, 4, 0)
```

## Iteration

```
for x in np.nditer(a):
```

Iterating over a with a prescribed  
broadcasting shape

Iterating over a in a row-major fashion

Iterating over a in a column-major fashion

```
for(auto it=a.begin(); it!=a.end(); ++it)
```

```
a.begin({3, 4})
a.end({3, 4})
```

```
a.begin<layout_type::row_major>()
a.begin<layout_type::row_major>()
```

```
a.begin<layout_type::column_major>()
a.end<layout_type::column_major>()
```



## NumPy

```
def diffusion(u, tempU, iterNum):  
    """  
    Apply Numpy matrix for the Forward-Euler Approximation  
    """  
    mu = .1  
    for n in range(iterNum):  
        tempU[1:-1, 1:-1] = u[1:-1, 1:-1] + mu * (  
            u[2:, 1:-1] - 2 * u[1:-1, 1:-1] + u[0:-2, 1:-1] +  
            u[1:-1, 2:] - 2 * u[1:-1, 1:-1] + u[1:-1, 0:-2])  
        u[:, :] = tempU[:, :]  
        tempU[:, :] = 0.0
```

## xtensor

```
auto diffusion(tensor<float, 2>& u, tensor<float, 2>& tempU, int iterNum)  
{  
    // Apply Numpy matrix for the Forward-Euler Approximation  
    float mu = 0.1;  
    for (int n = 0; n < iterNum; ++n)  
    {  
        view(tempU, _r|1|-1|, _r|1|-1|) = view(u, _r|1|-1|, _r|1|-1|) + mu * (  
            view(u, _r|2|_|, _r|1|-1|) - 2 * view(u, _r|1|-1|, _r|1|-1|) + view(u, _r|0|-2|, _r|1|-1|) +  
            view(u, _r|1|-1|, _r|2|_|) - 2 * view(u, _r|1|-1|, _r|1|-1|) + view(u, _r|1|-1|, _r|0|-2|)  
        );  
        u = tempU;  
        tempU.fill(0.f);  
    }  
}
```

# Xtensor is *fast*

Xtensor is continuously benchmarked for basic operation and a collection of real-life examples taken from the Pythran benchmarks.

Xtensor makes explicit use of the SIMD instructions available on the target platform for optimal performances. The xsimd project underlying the acceleration is its own project adopted beyond xtensor, and offers a new alternative to boost.simd.



Xtensor makes use of template expressions to implement Lazy Evaluation, unrolling most of the loops and preventing allocation of temporary variables.

# Xtensor performance

	Pairwise distance	Log likelihood	laplacian	cumsum	arc_distance	diffusion	Reverse cumsum
Python   NumPy	20020	4504	6032	3068	1080	20398	2522
xtensor	2934	2398	2210	1560	421	4949	13883
pythran	633	931	2964	1519	219	6179	2481
numba	2702	4005	8544	1476	1073	11764	1338
pypy	21991	4751	6001	2998	1150	22944	2519

A Sample from the Pythran Benchmark Suite

**Initialize a 2-D array and compute the sum of one of its rows and a 1-D array**

```
In [1]: #include <iostream>
#include "xtensor/xarray.hpp"
#include "xtensor/xio.hpp"
#include "xtensor/xview.hpp"
```

```
In [2]: xt::xarray<double> arr1
        {{1.0, 2.0, 3.0},
         {2.0, 5.0, 7.0},
         {2.0, 5.0, 7.0}};

xt::xarray<double> arr2
        {5.0, 6.0, 7.0};

std::cout << xt::view(arr1, 1) + arr2;;

{ 7., 11., 14.}
```

**Initialize a 1-D array and reshape it inplace**

```
In [3]: #include <iostream>
#include "xtensor/xarray.hpp"
#include "xtensor/xio.hpp"
```

```
In [4]: xt::xarray<int> arr
        {1, 2, 3, 4, 5, 6, 7, 8, 9};

arr.reshape({3, 3});

std::cout << arr;

{{1, 2, 3},
 {4, 5, 6},
 {7, 8, 9}}
```

**Broadcasting the `xt::pow` universal functions**

```
In [ ]: #include <iostream>
#include "xtensor/xarray.hpp"
#include "xtensor/xmath.hpp"
#include "xtensor/xio.hpp"
```

```
In [ ]: xt::xarray<double> arr3
        {1.0, 2.0, 3.0};

xt::xarray<unsigned int> arr4
        {4, 5, 6, 7};
```

# Trying xtensor online

With the new C++ Kernel in a Jupyter notebook

<https://mybinder.org/v2/gh/QuantStack/xtensor/0.16.4>



# What is coming?



Ongoing performance improvements:

- SIMD acceleration for complex numbers
- SIMD acceleration for `char` operations.



A C++ data frame. With a data model similar to that of pydata/xarray.

- A collection of N-D variables with labelled dimensions.
- Each variable is backed by an xtensor expression.


Just like xtensor, xframe is an ***expression system***, lazily evaluated.

# QuantStack

Scientific Computing



Sylvain Corlay



Wolf Vollprecht



Johan Mabil

# Resource

## CORE:

<https://github.com/QuantStack/xtl/>  
<https://github.com/QuantStack/xsimd/>  
<https://github.com/QuantStack/xtensor/>

## LANGUAGE BINDINGS

<https://github.com/QuantStack/xtensor-python/>  
<https://github.com/QuantStack/xtensor-r/>  
<https://github.com/QuantStack/xtensor-julia/>

## EXTENSIONS

<https://github.com/QuantStack/xtensor-io/>  
<https://github.com/QuantStack/xtensor-blas/>  
<https://github.com/egpbos/xtensor-fftw/>  
[https://github.com/wolfv/xtensor\\_ros/](https://github.com/wolfv/xtensor_ros/)

## XFRAME

<https://github.com/xframe>

## CORE:

<https://xtl.readthedocs.io>  
<https://xsimd.readthedocs.io>  
<https://xtensor.readthedocs.io>

## LANGUAGE BINDINGS

<https://xtensor-python.readthedocs.io>  
<https://xtensor-r.readthedocs.io>  
<https://xtensor-julia.readthedocs.io>

## EXTENSIONS

<https://xtensor-io.readthedocs.io>  
<https://xtensor-blas.readthedocs.io>